

## **ex1-JIT**

### **Step 1**

Java -XX:+PrintCompilation demo.MainV1 500

### **Step 2**

Java -XX:+UnlockDiagnosticVMOptions -XX:+LogCompilation  
demo.MainV1 7000

Note : hotspot\_pid19000.log file is created.

Search for random() and see the compiler C1 & C2

### **Step: 3**

Arguments-> Program Arguments:7000

VM arguments: -XX:+PrintCodeCache

### **Step: 4**

Tuning the code cache

VM Warning : Codecache is full. Compiler has been disabled

Set the below options :

#### Codecache Size Options

Option	Default	Description
InitialCodeCacheSize	160K (varies)	Initial code cache size (in bytes)
ReservedCodeCacheSize	32M/48M	Reserved code cache size (in bytes) - maximum code cache size
CodeCacheExpansionSize	32K/64K	Code cache expansion size (in bytes)

Now try with :

VM arguments: -XX:ReservedCodeCacheSize=28m -XX:+PrintCodeCache

#### **jvm compiler flags: -server, -client**

##### **Step1:**

Arguments-> Program Arguments: 10000

VM arguments -> -server -XX:+PrintCompilation

##### **Step2:**

Arguments -> Program Arguments: 10000

VM arguments -> -client -XX:+PrintCompilation

##### **Note:**

Disable TieredCompilation

VM arguments -> -XX:-TieredCompilation -XX:+PrintCompilation

#### **use jconsole to monitor the code cache**

note: jconsole process id is not showing, give write access to the below folder:

C:\Users\venkat\AppData\Local\Temp\hsperfdata\_venkat

> Memory Pool (Code Cache)

## **JVM Versions**

32-bit : Faster if heap < 3GB, Max heap size=4GB, Client Compiler only

64-bit : it would be faster if using long/double, heap > 4GB, Client & Server Compiler

## **Check the default java final flags**

```
java -XX:+PrintFlagsFinal
```

## **Check for the application process-id**

Check the CILCompilerCount value

[CILCompilerCount to two tells the JVM to use number of threads for compiler]

```
jinfo -flag CILCompilerCount <process-id>
```

## **Now change the count and check the output:**

VM arguments: -XX:CILCompilerCount=6 -XX:+PrintCompilation

## **CompileThreshold**

[Number of interpreted method invocations before (re-)compiling]

```
jinfo -flag CompileThreshold <process-id>
```

Now, try with the lower threshold value

```
-XX:CILCompilerCount=6 -XX:CompileThreshold=1000 -XX:+PrintCompilation
```

## Reading the compiler's mind

The -XX:+LogCompilation flag produces a low-level XML file about compiler and runtime decisions

-XX:+UnlockDiagnosticVMOptions -XX:+LogCompilation -  
XX:+PrintInlining -XX:+PrintCompilation

## Print Assembly:

-XX:+UnlockDiagnosticVMOptions -XX:+PrintAssembly

## Exception:

"Could not load hsdls-amd64.dll; library not loadable; PrintAssembly is disabled".

## Solution:

Download hsdls-1.1.1-win32-amd64.zip file from <http://fcml-lib.com/download.html> and copy hsdls-amd64.dll file in jdk bin folder.

(or) copy from this example lib folder

**Ref** : x86\_64 Assembly to understand generate code

## ex2.2-GC-Test

## GC Algorithms Benchmarking:

-Xms1g -Xmx1g -Xloggc:gc\_parallel2g.log -XX:+UseParallelGC -  
XX:+PrintGCDetails -XX:+PrintGCDateStamps

-Xms4g -Xmx4g -Xloggc:gc\_parallel4g.log -XX:+UseParallelGC -  
XX:+PrintGCDetails -XX:+PrintGCDateStamps

-Xms2g -Xmx2g -Xloggc:gc\_cms2g.log -XX:+UseConcMarkSweepGC -  
XX:+PrintGCDetails -XX:+PrintGCDateStamps

-Xms4g -Xmx4g -Xloggc:gc\_cms4g.log -XX:+UseConcMarkSweepGC -  
XX:+PrintGCDetails -XX:+PrintGCDateStamps

-Xms1g -Xmx1g -Xloggc:gc\_g1gc2g.log -XX:+UseG1GC -  
XX:+PrintGCDetails -XX:+PrintGCDateStamps

-Xms4g -Xmx4g -Xloggc:gc\_g1gc4g.log -XX:+UseG1GC -  
XX:+PrintGCDetails -XX:+PrintGCDateStamps

### **G1 GC With Pause time goal:**

-XX:MaxGCPauseMillis=200

## **ex2.3-GC-heap-sizes**

### **Monitor with VisualVM**

Run with :

```
java -Xmx20m -verbosegc Main  
Check the default value for UseAdaptiveSizePolicy  
jinfo -flag UseAdaptiveSizePolicy <process-id>
```

default : -XX:+UseAdaptiveSizePolicy

### **Set the ratios for Eden and Old generations**

jinfo -flag NewRatio <process-id>

-XX:NewRatio=n

[if n is 2 means; old space is two times bigger than eden space]

Note : By setting the NewRatio, UseAdaptiveSizePolicy will be disabled

### **Try with**

java -Xmx20m -XX:NewRatio=1 Main

-XX:SurvivorRatio=n

jinfo -flag SurvivorRatio <process-id>

### **Run with :**

java -Xmx20m -XX:SurvivorRatio=5

-XX:MaxTenuringThreshold=n

jinfo -flag MaxTenuringThreshold <process-id>

## **ex3.1-HeapDump-MAT**

1. Run the demo with -Xmx10m and -Xmx20m

2. Now replace the below code in CustomerManager, once again run the demo.

```
public Optional<Customer> getNextCustomer() {  
    synchronized(customers)  
    { if(customers.size()>0)  
        return Optional.of(customers.remove(0));  
    }  
    return Optional.empty(); }
```

### **Eclipse MAT (Memory analyzer Tool)**

File -> Acquire Heap dump -> Select the java process

Getting Started -> Leak Suspects Report

-> Problem Suspect 1-> Select Details

Analyse the "Shortest Paths To the Accumulation Point"

### **ex3.2 Memory Leak JMX**

1. Select ConsumerHeap.java -> Run As

-> Run configuration

-> set vm args

-XX:+HeapDumpOnOutOfMemoryError -Xms32m -Xmx32m

Error will be thrown :

java.lang.OutOfMemoryError: Java heap space

Dumping heap to java\_pid6664.hprof ...

Heap dump file created [1262470435 bytes in 5.784 secs]

Solution :

Analyse the error with "ECLIPSE MEMORY ANALYZER"

### **ex4.1-StringPool-Performance**

The String pool is implemented as a fixed capacity HashMap with each bucket containing a list of strings with the same hashCode.

#### **From command-prompt:**

```
java -XX:+PrintStringTableStatistics
```

Note : These Strings are loaded by the core java

#### **Try with application:**

```
java -XX:+PrintStringTableStatistics MainTest2
```

#### **Now, Try with:**

```
java -XX:+PrintStringTableStatistics -XX:StringTableSize=120121
```

Note : compare the Elapsed Time

#### **Tuning:**

```
java -XX:+PrintFlagsFinal MainTest2
```

Check the InitialHeapSize & MaxHeapSize



```
java -XX:+PrintStringTableStatistics -XX:StringTableSize=120121 -  
XX:MaxHeapSize=600m MainTest2
```

Note : Out of memory exception comes

### **Now, try with**

```
java -XX:+PrintStringTableStatistics -XX:StringTableSize=120121  
-XX:InitialHeapSize=1g  
compare the Elapsed Time
```

### **Shortcut Flags**

```
-XX:InitialHeapSize: -Xms  
-XX:MaxHeapSize: -Xmx
```

### **ex4.2-StringDeduplication**

4th execution of StringDeduplication. It took 35ms and edited 2,20,034 strings.

All were classified as "new", which means they have never been analyzed.

In the above example, all strings were deduplicated, which left a total of 8.5MB of memory.

"Total Exec" information is for summary of all invocations

```
-XX:+UseStringDeduplication -XX:+PrintStringDeduplicationStatistics
```

Enables the deduplication of Java Strings. This command-line option and feature was introduced in JDK 8u20. String deduplication is disabled by default.

`-XX:+PrintStringDeduplicationStatistics`

Enables the printing of String deduplication statistics. The default value is disabled. This command-line option can be very helpful when you want to know if enabling deduplication will result in a significant savings in the amount of space

in use by String objects in the Java heap. Hence, enabling this command-line option provides data that justifies whether there may be value in enabling

`-XX:+UseStringDeduplication.`

`-XX:StringDeduplicationAgeThreshold`

Sets the String object age threshold when a String object is considered a candidate for deduplication. The default value is 3.

More specifically, a String becomes a candidate for deduplication once a String object has been promoted to a G1 old region, or its age is higher than the deduplication age threshold.

### **ex5.1 Collections**

#### **clear() vs removeAll()**

the purpose of `clear()` and `removeAll(Collection c)` are different in API, `clear()` method is meant to reset a Collection by removing all elements, while `removeAll(Collection c)` only removes elements which are present in supplied collection.

This method is not designed to remove all elements from a Collection.

So, if your intention is to delete all elements from a Collection, then use `clear()`, while if you want to remove only some elements, which are present in another Collection,

e.g. list of closed orders, then use `removeAll()` method .

### **Clear() vs null**

Use `Clear` if we likely want to repopulate that list at some point. If it's a single use list and we are never going to use it again , then consider setting to null.

Note : There is a popular and widely accepted advice in the world of software development that says:

We should always return an empty list instead of null!

There are two considerable advantages:

We eliminate the risk of a null pointer error (i.e. `NullPointerException` in C#, `NullPointerException` in Java, etc.)

We don't have to check for null in client code - your code becomes shorter, more readable and easier to maintain

Contains in a linked list is proportional to the number of entries.

A `HashSet` will, on average, do the trick in constant time.

## **ex5.2 soft weak reference**

### **Soft References' Use Cases**

Soft references can be used for implementing memory-sensitive caches where memory management is a very important factor.

A cache can, for example, prevent its most recently used entries from being discarded by keeping strong referents to those entries, leaving the remaining entries to be discarded at the discretion of the Garbage Collector.

### **Weak References' Use Cases**

A weakly referenced object is cleared by the Garbage Collector when it's weakly reachable.

Weak references are most often used to implement canonicalizing mappings. A mapping is called canonicalized if it holds only one instance of a particular value. Rather than creating a new object, it looks up the existing one in the mapping and uses it.

### **ex5.3 phantom reference**

Phantom references have two major differences from soft and weak references.

We can't get a referent of a phantom reference. The referent is never accessible directly through the API and this is why we need a reference queue to work with this type of references.

The Garbage Collector adds a phantom reference to a reference queue after the finalize method of its referent is executed. It implies that the instance is still in the memory

### **Use Cases**

There're two common use-cases they are used for.

The first technique is to determine when an object was removed from the memory which helps to schedule memory-sensitive tasks. For example, we can wait for a large object to be removed before loading another one.

The second practice is to avoid using the finalize method and improve the finalization process.

### **lab-jmc**

#### **Customize the overview tab**

Select Dashboard - > + -> \*eden -> select used

#### **New Graph**

Select Historical Data Settings -> +

new chart is created

select the chart

Dashboard -> + -> Select Attributes -> \*sur -> PeakUsage -> Used

#### **MBean Browser**

Search \*eden -> select PeakUsage

Right click used -> visualize-> Chart

System

Memory

Threads

## **Diagnostic Commands**

Select GC.heap\_info -> Execute